

Predicting Postprandial Glycemic Response from Meal Composition and Subject Phenotype

A leave-one-subject-out study on the CGMacros cohort, with external validation on BIG IDEAs

Research module of the MetaNutri project. This is an independent, reproducible analysis — not a clinical tool, and not wired into the MetaNutri application. Datasets: CGMacros (PhysioNet, DOI 10.13026/3z8q-x658), licensed **CC BY-NC-SA 4.0**, and BIG IDEAs (PhysioNet, DOI 10.13026/zthx-5212), licensed **ODC-By 1.0**; see §9 for attribution.

Abstract

Postprandial glucose response (PPGR) varies several-fold between individuals eating the same meal, which makes personalised nutrition attractive but also hard to model. We reproduce and extend the CGMacros baseline on a meal-level table built from 45 subjects and 1,699 continuous-glucose-monitoring (CGM) annotated meals. Using strictly subject-wise **leave-one-subject-out (LOPO)** cross-validation, we predict three response targets — 2-hour incremental area under the curve (iAUC), total area under the curve (AUC), and peak glucose rise — from meal macronutrients plus subject blood-panel and anthropometric features. On the breakfast-only subset our XGBoost model reaches **$r = 0.890$ for AUC** and **$r = 0.655$ for iAUC**, matching the published baseline (≈ 0.89 / ≈ 0.64), which validates the pipeline. Extending the same features and model to all meal types (1,557 meals) yields **$r = 0.838$ (AUC)** and **$r = 0.451$ (iAUC)**. A mean predictor and two macronutrient-only Ridge regressions are far weaker (best linear $r = 0.255$ for breakfast iAUC), showing that most of the signal is non-linear or subject-specific. We then test cross-cohort generalisation by freezing the CGMacros model and applying it, unchanged, to the independent **BIG IDEAs** cohort (16 subjects, 656 meals, different CGM device and free-living self-reported food logs). AUC transfers partially (**$r = 0.569$**), but iAUC and peak rise collapse to $r \approx 0.22$ — and a simple carbohydrate-only regression transfers better ($r = 0.36$) than the tree ensemble, a cautionary result about non-linear models that fit cohort-specific structure. We report limitations candidly: small cohorts, pooled-observation metrics that mix between- and within-subject variance, and a large domain gap between the two studies. Code, figures and fixed-seed experiments are included for full reproducibility.

1. Problem

Two people can eat an identical meal and show very different glucose excursions. Capturing that heterogeneity is the core promise of personalised nutrition, and also its core difficulty: the response depends on meal composition, on the person's metabolic state, and on interactions between the two.

We ask a deliberately narrow, testable question: **given a meal's macronutrients and a set of subject-level clinical features, how well can we predict the 2-hour postprandial glucose response, when the model is evaluated on people it has never seen?** The emphasis on unseen subjects is the point. A model that memorises a person's typical response will look excellent under a random split and fail in deployment, so every number in this report comes from LOPO cross-validation.

The intended use of this work is methodological: it is a writing and engineering sample that demonstrates a correct evaluation protocol on real data. It is not a medical device and makes no diagnostic claim.

2. Data

2.1 Source and cohort

CGMacros provides CGM traces with time-stamped meal annotations and macronutrient estimates, plus a subject-level `bio.csv` with laboratory measurements. The public release contains 45 subjects. Table 1 summarises the analysis cohort after filtering.

Table 1 — Analysis cohort.

Property	Value
Subjects	45
Meals with iAUC > 0	1,557 (of 1,699 extracted)
Meals per subject (min / median / max)	17 / 35 / 61
Meal types	breakfast 423, lunch 380, dinner 451, snack 303
Age (min / median / max)	18 / 51 / 69
BMI (min / median / max)	20.7 / 30.0 / 49.1
HbA1c % (min / median / max)	4.6 / 5.9 / 8.5
Glycemic status (healthy / prediabetes / T2D)	573 / 576 / 408 meals
Sex (M / F)	16 / 29

The cohort deliberately spans healthy, prediabetic and type-2-diabetic ranges, so the model must handle a wide dynamic range rather than a single narrow population. The dataset is small by machine-learning standards: ~1.5k meals and 45 independent subjects. This shapes every modelling choice below.

2.2 From CGM traces to a meal-level table

For each meal annotation we take the 2-hour postprandial window, sample the glucose series every 15 minutes (9 points), and compute three targets:

- **iAUC** — positive incremental area above the meal-time baseline, using trapezoidal integration with correct handling of curve crossings;
- **AUC** — total trapezoidal area over the window;
- **peak rise** — maximum glucose elevation above baseline.

Features are the four macronutrients (carbohydrate, protein, fat, fiber), the meal-time baseline glucose, and 14 subject-level variables including age, sex, BMI, HbA1c, fasting glucose, insulin, HOMA-IR and a lipid panel (triglycerides, total cholesterol, HDL, non-HDL, LDL, VLDL, cholesterol/HDL ratio). Subject features are joined **on subject id**, not positionally.

Two data-quality issues were handled explicitly rather than silently. First, column names are inconsistent across the 45 per-subject files (trailing spaces, and one file lacks Amount Consumed); we normalise names and access optional columns defensively. Second, meal-type labels are inconsistent (Snacks, snack 1, Breakfast, ...); we collapse them to breakfast / lunch / dinner / snack. Only one value is missing across the final feature matrix (a single fiber_g), imputed with the **training-fold median**.

Figure 3 shows the resulting data: iAUC distributions differ markedly by meal type, and carbohydrate content correlates only weakly with iAUC ($r \approx 0.2$), foreshadowing the modest ceiling of any meal-composition-only model.




Figure 3 — Data overview. iAUC by meal type; carbohydrate content versus iAUC; meal counts per type.

2.3 External cohort: BIG IDEAS

To test whether the model generalises beyond CGMacros, we use the **BIG IDEAs Lab Glycemic Variability and Wearable Device** dataset as an independent external cohort. It contributes 16 subjects with Dexcom G6 CGM (5-minute) and a free-living food log. Three differences from CGMacros matter for the analysis:

1. **Meals are not annotated.** The release provides a per-item food log; we group items logged within 20 minutes into one eating event and sum their macronutrients, matching the CGMacros target window (2 h, 9-point sampling).
2. **A narrower feature set.** BIG IDEAs does not record the CGMacros blood panel, so only the five shared features (carbohydrate, protein, meal-time baseline glucose, sex, HbA1c) are used for the primary comparison — complete for all 16 subjects. A seven-feature variant adding fat and fibre covers 13 subjects (three subjects log fat/fibre sparsely or not at all) and is reported as a sensitivity check.
3. **Food-log dates are misaligned in the version we could obtain.** The dataset's own 1.1.3 release notes state "Updated misaligned food log dates", confirming the defect. Because 1.1.3 was not retrievable (HTTP 403) and the open mirror carries only 1.1.2, we use 1.1.2 and repair the dates ourselves: for each subject we search integer day offsets and pick the shift that maximises the median post-meal glucose rise, subject to retaining ≥ 90 % of the maximum achievable event coverage. This yields no shift for 12

subjects and a roughly five-month shift for four (007, 013, 015, 016), consistent with the reported defect. As an independent check, the study's standardised-breakfast days show a higher morning glucose rise than other days after correction (median $\Delta \approx +14$ mg/dL).

After correction, the BIG IDEAs table contains **656 meals across 16 subjects**.

3. Method

3.1 Evaluation protocol

We use **leave-one-subject-out cross-validation**. In each of the 45 folds, all meals from one subject are held out; the model is fit on the remaining 44 subjects and predicts the held-out subject's meals. Standardisation and median imputation are fit on the training fold only, so no information from the test subject leaks into preprocessing. Metrics are computed on the pooled out-of-fold predictions.

This matters quantitatively. Random meal-level splits let the same person appear in train and test, and individual response baselines are highly subject-specific, so random splits substantially overstate performance. Any single-speaker result that is not subject-wise should be read with that optimism in mind.

3.2 Models

We compare four models, deliberately spanning a naive floor and a strong tree-based learner:

1. **Mean predictor** — predicts the training-fold mean. This is the honest floor: anything that cannot beat it has learned nothing.
2. **Carb-only Ridge** — linear regression on carbohydrate.
3. **Energy-only Ridge** — linear regression on carbohydrate, protein and fat.
4. **XGBoost** — gradient-boosted trees with the official CGMacros baseline hyper-parameters: `max_depth = 1`, `n_estimators = 80`, `learning_rate = 0.2`, `reg_alpha = 1.0`.

Ridge, not ordinary least squares, is used for the linear models (`alpha = 1.0`). With 19 strongly collinear features and subjects carrying extreme values, unregularised OLS extrapolates catastrophically on held-out subjects (we observed R^2 below $-1,000$); a small L2 penalty makes the baseline comparison meaningful rather than a numerical artefact. The tree depth of 1 makes each boosted round a decision stump, which is what keeps the model stable in this small-sample regime.

3.3 Metrics

Our primary metric is the **Pearson correlation r** between observed and predicted values under LOPO, which is what the CGMacros baseline reports and therefore what makes replication directly comparable. We additionally report Spearman ρ , R^2 , RMSE and MAE. R^2 and RMSE are on the original, untransformed scale (iAUC in mg/dL·min), and are more sensitive to the long right tail than r is; we therefore treat r as the headline and the error metrics as context. Prediction intervals are not reported because the tree ensemble does not natively produce calibrated uncertainty for unseen subjects.

4. Results

4.1 Breakfast: replication of the published baseline

Restricting to breakfast (423 meals, 45 subjects) reproduces the official protocol. Table 2 and Figure 2 show the outcome.

Table 2 — Breakfast subset (n = 423, LOPO).

Model	AUC r	iAUC r	Peak r	iAUC R ²	iAUC RMSE
Mean predictor	−0.924	−0.820	−0.826	−0.03	3,474
Carb-only Ridge	−0.014	0.107	0.079	0.00	3,414
Energy-only Ridge	0.100	0.255	0.247	0.06	3,313
XGBoost (official)	0.890	0.655	0.632	0.43	2,589

Our XGBoost result of **r = 0.890 (AUC)** and **r = 0.655 (iAUC)** aligns with the published ≈ 0.89 and ≈ 0.64 , which is the main validation that our dataset construction, target definitions and evaluation split are correct. The agreement is close enough that we treat the pipeline as a faithful reproduction, not merely a similar-looking re-run.

The linear baselines expose the structure of the problem. Carbohydrate alone barely correlates with iAUC ($r = 0.107$); adding protein and fat helps ($r = 0.255$), consistent with fat and protein modulating gastric emptying and the glycemic curve. But even the best linear model captures only $\sim 6\%$ of iAUC variance, versus $\sim 43\%$ for XGBoost. The gap is the empirical case for a non-linear model here.

The mean predictor's strongly negative correlation is a useful diagnostic rather than a real result. Because each fold's prediction is the mean of the remaining subjects, a subject with an unusually high response pulls that mean down when held out — inducing an anti-correlation. It confirms that between-subject differences dominate the pooled variance, and it is exactly why a subject-wise split is mandatory.

4.2 All meals: an extension

We then applied the identical features and model to all meal types (1,557 meals), removing the breakfast restriction. Table 3 and Figure 1 summarise.

Table 3 — All-meals subset (n = 1,557, LOPO).

Model	AUC r	iAUC r	Peak r	iAUC R ²	iAUC RMSE
Mean predictor	−0.780	−0.553	−0.552	−0.01	2,735
Carb-only Ridge	−0.032	0.147	0.150	0.02	2,692
Energy-only Ridge	−0.035	0.143	0.146	0.01	2,697
XGBoost (official)	0.838	0.451	0.542	0.20	2,431

XGBoost LOPO predictions (each point = one held-out meal)




Figure 1 — Predicted versus observed under LOPO. Each point is one held-out meal. Left column: breakfast only; right column: all meals.

AUC remains well predicted ($r = 0.838$) across meal types. iAUC drops to $r = 0.451$ and peak rise to $r = 0.542$. This is the expected direction: breakfasts are the most standardised meal, while lunches, dinners and snacks vary far more in composition, timing and context, and the pooled iAUC distribution becomes heavier-tailed. The model generalises across meal types, but with a clear and honest degradation.




Figure 2 — Model comparison. Out-of-fold Pearson r by model, target and subset. XGBoost dominates everywhere; the linear baselines are close to zero, and the mean predictor is negative by construction.

4.3 What the model uses

Figure 4 ranks XGBoost gain importance for iAUC. The dominant features are meal carbohydrate, the meal-time **baseline glucose**, and the subject's glycemic state (HbA1c, fasting glucose). This ordering is physiologically sensible and reassuring: the model leans on the current glucose level and long-run glycemic control, not on incidental lipid variables. Fiber and the lipid panel contribute little once these are present.




Figure 4 — Feature importance for iAUC. Descriptive fit on all available data (not LOPO); shown for interpretation, never used for model selection.

4.4 External validation on BIG IDEAs

The frozen CGMacros model (same 80 depth-1 trees, same five shared features) is applied to the 656 BIG IDEAs meals without any retraining or refitting of the preprocessing beyond the CGMacros-fitted scaler. We report three quantities side by side: the model's **internal** LOPO score on each cohort, and its **external transfer** score. Table 4 and Figure 5 show the outcome.

Table 4 — XGBoost Pearson r, 5 shared features.

Cohort / protocol	iAUC r	AUC r	Peak r
CGMacros (LOPO, internal)	0.487	0.816	0.508
BIG IDEAs (LOPO, internal)	0.463	0.625	0.483
BIG IDEAs (external transfer)	0.227	0.569	0.223
BIG IDEAs (external, carb-only Ridge)	0.362	0.288	0.357
BIG IDEAs (external, mean predictor)	— ($R^2 < 0$)	— ($R^2 \approx 0$)	— ($R^2 < 0$)

Two things stand out. First, **AUC transfers far better than iAUC**. A correlation of 0.569 for AUC across studies that differ in participants, device and meal annotation is a genuine, if modest, positive result: part of the AUC signal is a portable, physiology-level relationship (glucose level and long-run control). Second, **iAUC and peak rise do not transfer**: r falls from 0.487 and 0.508 internally to 0.227 and 0.223 externally, and R^2 is near zero. The model has learned CGMacros-specific structure for the excursion that does not generalise to free-living meals in another cohort.

The internal BIG IDEAs LOPO column is the important control. It shows that BIG IDEAs itself is not unlearnable — a model trained and tested within that cohort reaches $r = 0.463$ for iAUC. The drop to 0.227 is therefore attributable to **domain shift across cohorts**, not to noise in the external data.

The most uncomfortable result is the baseline comparison. On external iAUC and peak rise, the **carbohydrate-only Ridge regression transfers better** ($r = 0.362$ and 0.357) than the XGBoost model (0.227 and 0.223). The tree ensemble, which wins decisively inside CGMacros, appears to have fit cohort-specific interactions that do not survive the transfer, whereas the single, near-universal "more carbohydrate $\rightarrow$ larger excursion" relationship is more robust. The seven-feature variant adding fat and fibre changes little (external AUC $r = 0.581$, iAUC $r = 0.272$, $n = 558$). We treat this as a cautionary finding rather than a failure: it is exactly the kind of result that an external test is meant to expose, and it would be invisible under any within-cohort evaluation.




Figure 5 — External validation. Left: XGBoost Pearson r by cohort and target under the five shared features; the red bars are the frozen-model transfer. Right: predicted versus observed for the transferred model — the AUC cloud follows the diagonal loosely, while iAUC and peak rise are nearly flat.

5. Discussion

Four findings are worth stating plainly.

First, **the published baseline reproduces**, and the extension is credible. Matching $r \approx 0.89$ (AUC) and $r \approx 0.64$ (iAUC) on breakfast means subsequent numbers sit on a validated pipeline rather than an idiosyncratic one.

Second, **non-linearity earns its place**. The best linear model reaches $r = 0.255$ on breakfast iAUC, whereas the boosted stumps reach 0.655. With only 80 depth-1 trees, XGBoost is still a small model; the gain comes from modelling thresholds and interactions (for example, the effect of carbohydrate differing by baseline glucose), which a single linear coefficient cannot express.

Third, **AUC is easier than iAUC**, consistently. AUC is dominated by the subject's overall glucose level, which the baseline and HbA1c features capture well. iAUC subtracts that baseline and isolates the excursion, the part that depends on the meal and on individual handling of it — and that is harder. This distinction is often blurred in PPGR reporting, where a high AUC correlation can be mistaken for good meal-response modelling.

Fourth, **external validation changes the story in a useful way**. Inside CGMacros, XGBoost dominates the linear baselines on every target. Across cohorts, the ranking partly reverses: the tree model still transfers for AUC ($r = 0.569$) but loses to a one-variable carbohydrate regression on iAUC and peak rise. The lesson is not that gradient boosting is wrong — it is that a model whose advantage comes from cohort-specific interactions carries a hidden generalisation risk, and only an external test reveals it. We would rather report this than present a single-cohort number that flatters the method.

6. Limitations

We state these explicitly; they bound every claim above.

- **Small cohort.** 45 subjects cannot support deep models or fine-grained subgroup analysis. Neural approaches that need tens of thousands of samples are not justified here, which is why we use gradient-boosted stumps.
- **Pooled metrics mix two variance sources.** Correlations are computed over pooled meals, so between-subject differences inflate them relative to a within-subject evaluation. A per-subject standardised variant would be stricter and would likely lower the numbers; we report the comparable metric but flag the effect.
- **External validation is narrow and imperfect.** We test one external cohort (BIG IDEAs, 16 subjects, 656 meals) with only five shared features, because that cohort lacks the CGMacros blood panel. The two studies differ simultaneously in participants, CGM device (Libre Pro vs Dexcom G6) and meal annotation (expert-annotated vs free-living self-report), so the drop in performance cannot be attributed to any single factor. The BIG IDEAs meals are also derived from a 1.1.2 food log whose dates required our own repair procedure, which is validated but not ground truth.
- **No uncertainty quantification.** Point predictions are reported without calibrated intervals; for a health-adjacent setting, that is a real gap.

- **Observational and non-causal.** Associations between features and response are not interventions. Nothing here implies that changing a macronutrient will cause the predicted change.
- **Data provenance caveats.** The dataset's internal `LICENSE.txt` is an empty file; the CC BY-NC-SA 4.0 terms are declared only on the PhysioNet web page. We treat the web page as authoritative and comply with attribution and share-alike.

7. Future work

1. **More external cohorts and domain adaptation**, building on the BIG IDEAs result. Re-calibration or a small amount of target-cohort data may recover the iAUC signal that the frozen model loses; testing a third cohort (e.g. a T2D population) would separate device effects from population effects.
2. **Within-subject evaluation** in addition to pooled metrics, to separate "who is this person" from "what did they eat".
3. **Uncertainty estimates** (quantile or conformal prediction) so a prediction comes with a defensible interval.
4. **Richer meal representations** — meal timing, prior-meal context, and activity — which the current feature set omits.
5. **An interactive demo** exposing per-meal predictions and feature attributions, built on this frozen, reproducible pipeline.

8. Reproducibility

All results are produced by a small, dependency-pinned pipeline (`research/`), independent of the MetaNutri platform:

```
src/download_data.py      # parallel-chunk download + SHA256 verification
src/build_dataset.py      # CGM → meal-level table (targets + features)
src/iauc.py               # iAUC / AUC / peak-rise definitions
src/evaluate.py           # LOPO protocol, imputation, baselines, XGBoost
src/experiment.py         # runs both subsets × three targets → results.csv
src/download_bigideas.py  # fetch the 33 small BIG IDEAs files (2.4 MB)
src/build_external.py     # BIG IDEAs → meal-level table (date-offset repair)
src/external_validate.py  # cross-cohort transfer → external_results.csv
src/make_figures.py       # figures 1–5
src/export_report_pdf.py  # this report → reports/technical_report.pdf
```

Pinned versions: Python 3.12, `numpy==2.5.3`, `pandas==3.0.6`, `scikit-learn==1.9.1`, `xgboost==3.4.1`, `scipy==1.18.1`, `matplotlib==3.11.2`. Random seed 42. Re-running `experiment.py` regenerates `experiments/results.csv`; `make_figures.py` regenerates the figures.

9. Data attribution and license

This work uses the **CGMacros** dataset:

- Source: PhysioNet — <https://physionet.org/content/cgmacros/1.0.0/>

- DOI: [10.13026/3z8q-x658](https://doi.org/10.13026/3z8q-x658)
- Reference: Das et al., Scientific Data 12, 1557 (2025).
- License: **Creative Commons Attribution-NonCommercial-ShareAlike 4.0 (CC BY-NC-SA 4.0)**.

Use is non-commercial (research and portfolio). This analysis is a derivative work and is distributed under the same license, with attribution to the original authors. Raw data is never redistributed in this repository.

External validation uses the **BIG IDEAs Lab Glycemic Variability and Wearable Device Data** dataset:

- Source: PhysioNet — <https://physionet.org/content/big-ideas-glycemic-wearable/1.1.2/>
- DOI (version 1.1.2, used here): [10.13026/zthx-5212](https://doi.org/10.13026/zthx-5212)
- Reference: Cho, P., Kim, J., Bent, B., & Dunn, J., BIG IDEAs Lab Glycemic Variability and Wearable Device Data, PhysioNet.
- License: **Open Data Commons Attribution License 1.0 (ODC-By 1.0)**, which permits reuse with attribution.